

What does the crowd believe? A hierarchical approach to estimating subjective beliefs from empirical data

Erin Bennett, Judith Degen, Michael Henry Tessler, Noah D. Goodman

{ebennett,jdegen,mhtessler,ngoodman}@stanford.edu

Department of Psychology, 450 Serra Mall

Stanford, CA 94305 USA

Michael Franke

michael.franke@uni-tuebingen.de

Department of Linguistics, Wilhelmstraße 19

72074 Tübingen, Germany

Each author's name should appear on a separate line, 11 point bold, and centered, with the author's email address in parentheses. Under each author's name list the author's affiliation and postal address in ordinary 10 point type.

who is an author? which order?

Abstract

Subjects' beliefs about everyday events are of theoretical interest on their own and also an important ingredient for model building—especially in Bayesian models of phenomena such as logical reasoning, future predictions, and language use. Here, we explore several recently used methods for measuring subjective beliefs about unidimensional properties. [ndg: say what kind of subjective beliefs /distributions we look at more clearly.] We use a hierarchical Bayesian data analysis model for inferring likely “population-level beliefs” as the central tendency of subjects' individual-level beliefs. Three different dependent measures—[ndg: name them?]
—are related to individual-level beliefs through likelihood “linking functions.” Our results suggest that the practice of using averaged normalized slider ratings for binned quantities is a practical and fairly good approximator of inferred population-level beliefs. Other dependent measures yield less reliable results. **Keywords:** subjective beliefs, hierarchical modeling, Bayesian data analysis, Bayesian cognitive models

Motivation

We cannot look into a person's head. But in making sense of observed behavior we readily ascribe beliefs and desires to fellow agents. This happens intuitively, in folk psychology, but also in science. Scientific ascription of latent mental states plays an important role in many explanations of higher-order cognition: decision making, planning, reasoning, and language use. Naturally, then, it becomes important to validate any method for eliciting mental states, to be used for explanatory purposes.

A family of models where this is particularly pressing are Bayesian models of cognition which seek to explain task behavior in a variety of domains as partially informed by what subjects believe about mundane events—their *prior beliefs*. Take interpretation of language. Empirical data on whether a statement like “That watch cost n dollars” is understood to convey speaker affect, rather than literal meaning, can be

explained well by a Bayesian model of utterance interpretation (Kao, Wu, Bergen, & Goodman, 2014). In this model a crucial role is played by an empirical measure of subjects' expectations about the likely or normal price of a watch. Other examples of domains in which empirically successful models have included some measure of subjects' belief include making future predictions (Griffiths & Tenenbaum, 2006), the strength of pragmatic enrichments (Degen, Tessler, & Goodman, 2015), the interpretation of vague quantifiers (Schöller & Franke, 2015), ...

fill me

Many methods of assessing subjective beliefs exist. One approach is to take actual frequencies as an approximation (Griffiths & Tenenbaum, 2006). This will not work when sufficient real-world data is unavailable—for instance, for beliefs about imaginary events. Even where available, real-world frequencies may deviate systematically from subjects' beliefs, and this can be problematic for the predictions of a model that relies on subjective beliefs.

Another approach is to experimentally elicit relevant beliefs. Many techniques for this exist, especially in the economics literature (Morgan & Henrion, 1990; Manski, 2004; Schlag, Tremewan, & van der Weele, online first; Andersen, Fountain, Harrison, & Rutström, 2014). The perhaps most prominent method uses so-called *scoring rules* (Savage, 1971; Schlag et al., online first). Roughly speaking, scoring rules are particular schemes of rewarding participants' answers in such a way as to make sure that responses are what we want them to be: the actual subjective probability of an event, the mean of a distribution, a credible interval etc. Scoring rules are highly faithful but require participants to become sufficiently familiar with the (usually probabilistic) payoff scheme. Scoring rules also must make assumptions about how much participants value certain probabilistic gambles and may therefore require elaborate means of data analysis (Andersen et al., 2014). Similar practical concerns apply to other methods as well, such as the use of *iterated learning* for prior estimation (Lewandowsky, Griffiths, & Kalish, 2009).

For common applications, many of these procedures may

be excessively complex or too demanding on both researcher and participant. It would be beneficial if there was an easily applicable technique of eliciting subjective beliefs that is good enough for testing the general predictions of models that rely on these measures for their predictions. A further complication is that it may be impractical to obtain, from the same participant, data on subjective beliefs and data on the task of interest, especially if exposure to one of these measures is likely to affect the other. This suggests that, for practical reasons, there is a demand of reliable-enough measures of what we will sloppily call “population-level beliefs”: estimates of the central tendency or average of subjective beliefs in a given population.

A very simple technique that has been used with apparent success is to have subjects adjust sliders in order to express their (relative) levels of support for values of uncertain contiguous quantities (Kao et al., 2014; Schöller & Franke, 2015). **MORE REFERENCES** Subjects are confronted with a short and simple description of a situation or an event, such as “John bought a watch,” and asked for their estimates of the watch’s price. To give their estimate, they are asked to adjust sliders, whose ends are labelled as “extremely likely” and “extremely unlikely,” for a number of bins that partition the range of plausible values (established by a pre-test). E.g., for the price of a watch bins could be **FILL ME**

[ndg: should we also talk about the other dependent measures? i think part of what’s interesting is comparing the different dependent measures to each other?]

[ndg: say more clearly that we focus on unidimensional distributions and why.]

- explain slider rating task
- our goal here is to validate whether “population-level beliefs” are approximated well enough by normalized average slider ratings
- use hierarchical Bayesian modeling
 - calculate likelihoods from subjective beliefs
 - each subjective belief is drawn from a population-level central tendency, which is a “belief” as well

Experiments

- why within-participant?
- what was the exact randomization / procedure?
- items are given in Table 1
- include bins as well?
- what were the lighting round bin comparisons?
 - 1 vs 2
 - 2 vs 6
 - 6 vs 11
 - 11 vs 14

1. X has just fetched himself a cup of **coffee** from the office vending machine. What do you think the temperature of his coffee is?
2. X **commuted** to work yesterday. How many minutes do you think she spent commuting yesterday?
3. X told a **joke** to N kids. How many of the kids do you think laughed?
4. X bought a **laptop**. How much do you think it cost?
5. X threw N **marbles** into a pool. How many of the marbles do you think sank?
6. X just went to the **movies** to see a blockbuster. How many minutes long do you think the movie was?
7. X watched **TV** last week. How many hours do you think he spent watching TV last week?
8. X bought a **watch**. How much do you think it cost?

Table 1: Experimental items

– 14 vs 15

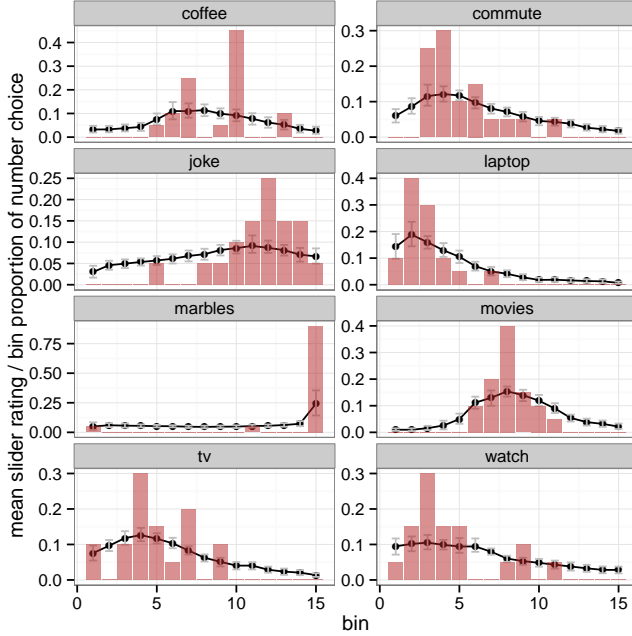
- results are in Figure 1

Model


The data we would like to explain are: (i) the normalized slider ratings $s_{ijk} \in [0; 1]$ of subject $i \in \{1, \dots, 20\}$ for item $j \in \{1, \dots, 8\}$ and bin $k \in \{1, \dots, 15\}$; (ii) the bins $n_{ij} \in \{1, \dots, 15\}$ in which subject i ’s number choice for item i was; and (iii) the binary choices $c_{ijl} \in \{0, 1\}$ of whether subject i selected the higher bin for item j in the bin comparison condition l . There two simplifications in need of commenting. In (i), we focus on slider ratings after normalizing for each subject, because we assume that slider adjustments reflect relative, not absolute estimates of subjective beliefs. In (ii), we focus on bin choices, not actual number choices, in order to avoid, as much as possible, considerations of salience of particular numbers, and also because otherwise data from items with smaller domains of plausible number would get more weight than data from items which allow for a wider set of number choices.

All three pieces of data are to be explained as functions of subjective beliefs P_{ij} , where P_{ij} is a probability vector of length 15 with P_{ijk} being subject i ’s belief about the relative likelihood of bin k for item j . Each P_{ij} defines a likelihood for our data, via appropriate link functions (to be spelled out presently). Variance in subjective beliefs is harnessed by a population-level hyper-prior with central tendency Q_j , i.e., a stochastic vector of length 15, which we call a “population-level belief” about item j . The structure of this model is pictured in Figure 2.

To fill the structure in Figure 2 with life, we need to spell out three parameterized link functions, one for each task type, and the relation between population-level belief Q_j to individual beliefs P_{ij} . Let’s start with the latter. For fixed Q_j , the idea is that P_{ij} are noise-perturbed variants scattered around



(a) slider ratings & number choices



(b) bin comparison

Figure 1: Results from experiments. Error bars show bootstrapped 95% confidence intervals.

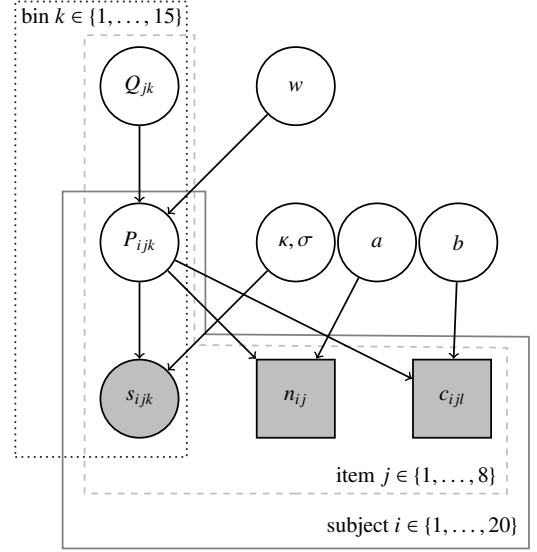


Figure 2: The data-generating model as a probabilistic graphical model, following conventions of Lee and Wagenmakers (2015). Shaded nodes are observed, white nodes are latent variables. Square nodes represent categorical, round nodes continuous variables. Boxes indicate scope of indices.

Q_j , with some parameter w to determine how much perturbation we should expect. To realize this, the model assumes that P_{ij} are distributed according to a Dirichlet distribution with weights given by wQ_j :

$$Q_j \sim \text{Dirichlet}(1, \dots, 1) \quad w \sim \text{Gamma}(2, 0.1)$$

$$P_{ij} \sim \text{Dirichlet}(wQ_j)$$

The higher w , the more likely it is that P_{ij} is “close” to Q_j .

The link function for the slider rating data uses a logit transformation to project observed slider ratings s_{ijk} and latent probabilities P_{ijk} , which are bound to lie between 0 and 1, to the reals. The likelihood of logit-transformed observation s_{ijk} is given by a Gaussian with standard deviation σ around the logit-transformed predictor P_{ijk} . On top of that, there is a parameter κ , the steepness of the logit transform of P_{ijk} , that allows response likelihoods to capture end-point affinity for $\kappa > 1$ (values of P_{ijk} close to 0 or 1 are likely mapped to 0 or 1) or end-point aversion for $\kappa < 0$ (values of P_{ijk} are likely to be realized as more median), with a prior that expects $\kappa = 1$.

$$\text{logit}(s_{ijk}) \sim \text{Norm}(\text{logit}(P_{ijk}, \kappa), \sigma)$$

$$\sigma \sim \text{Gamma}(0.0001, 0.0001) \quad \kappa \sim \text{Gamma}(5, 5)$$

The link function for number choice data treats each bin n_{ij} as a draw from a categorical distribution where the probability of bin k is proportional to $\exp(aP_{ij})$, i.e., a soft-max choice from P_{ij} . The higher parameter a , the more likely n_{ij} is the mode of P_{ij} . For $a \rightarrow 0$, all bins become equiprobable.

$$n_{ij} \sim \text{Categorical}(\exp(aP_{ij})) \quad a \sim \text{Gamma}(2, 1)$$

Finally, consider the link function for bin comparisons. We are after the likelihood with which subject i selects the higher bin for item j in bin comparison condition l . Suppose l is about comparing the lower bin b_l to the higher b_h . The perhaps most natural approach would be to link the likelihood of choosing b_h over b_l to the difference between P_{ijb_h} and P_{ijb_l} . This, however, does not appear to be what subjects were doing. For example, take the “marbles” case where we see almost unanimous choices of bin 11 over bin 6, despite seeing any pronounced difference in slider ratings for these bins. Indeed, a model that implements this idea failed to capture the relevant regularities in the data blatantly.

Another plausible link function that accommodates for this is to assume that what matters in bin comparisons is the distance to the mode of P_{ij} : soft-max prefer the bin that is closer to the mode of P_{ij} ; select random if both bins are equally far from the prototype.¹

$$c_{ijl} \sim \text{Bern}((1 + \exp(2b(1 - p_{ijl}^{\text{high}})))^{-1}) \quad b \sim \text{Gamma}(2, 1)$$

$$p_{ijl}^{\text{high}} = \begin{cases} 2 & \text{if mode}(P_{ij}) \text{ is closer to higher} \\ & \text{bin of } l \text{ than to lower bin} \\ 1 & \text{if equal distance} \\ 0 & \text{otherwise} \end{cases}$$

[ndg: wow, this is weird! you tried softmax ratio of bin probs?]

Inference

The model was implemented in JAGS (Plummer, 2003). 50,000 samples were obtained from two chains with a thinning rate of 2 after a burn-in of 100,000 that ensured convergence according to \hat{R} (Gelman & Rubin, 1992).

Of most theoretical interest are posterior estimates of Q_j . Figure 3 shows means of the posterior Q_j , with their 95% HDIs, alongside the averaged normalized slider ratings. The latter provide a very good approximation of the latent, inferred population-level beliefs. Inspection of posteriors of individual P_{ij} shows that there is ample variation between subjects. Nonetheless, the way the model suggest we should think about harnessing the individual P_{ij} s under a population-level central tendency is closely approximated by simply averaging over slider ratings. Although the match is clearly not perfect, it is good enough to maintain that the latter are a reasonable (enough) and practical way of assessing what the crowd believes despite individual differences.

[ndg: it'd be nice to see some posteriors for some of the (more meaningful) parameters. eg w, kappa, ...]

[ndg: what about the success of the other DMs? looks like give-n is not so hot. worth saying that if so.]

¹Notice that $(1 + \exp(2b(1 - p_{ijl}^{\text{high}})))^{-1} = \frac{\exp(bp^{\text{high}})}{\exp(bp^{\text{high}}) + \exp(bp^{\text{low}})}$ with $p_{ijl}^{\text{low}} = 2 - p_{ijl}^{\text{high}}$.





Figure 3: Means of posteriors over Q_j in black with gray area indicating 95% HDIs. Red lines give the average normalized slider ratings for comparison.

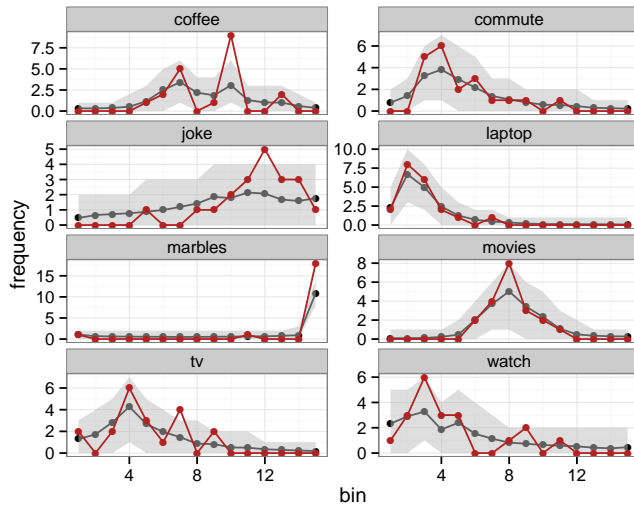
Model criticism

Inferences based on the model are only as reliable as the model itself is plausible. Model criticism is therefore important. Figure 4 shows posterior predictive checks at the population-aggregate level for all of our three task types. For the slider rating task, posterior predictions are spot-on. Some of the number estimation data is surprising even for the model trained on this very data. This could have various reasons: (i) the number data does not have a huge influence on the posterior likelihood, (ii) number choices may be influenced by saliency and/or roundness of numbers after all. Finally, there is one condition in the bin comparison task that the model definitely got wrong. This is the choice of what is more likely: that one or that none of 14 marbles thrown into a pool would sink. The model would predict that almost everybody should answer that it is more likely that one marble sank. But that is not what we observe. It may be that subjects revise beliefs about “normality” of the marbles, while holding on to an assumption that all marbles behave the same (Degen et al., 2015).


Posterior predictive checks indicate that the trained model captures patterns of answers at the aggregate population level well. To have a more fine-grained measure of model fit, we also looked at posterior predictive p -values at the level of subjects and items. [ndg: need to say what this means – i’m not even sure that i know!] Unfortunately, only the slider rating task, provides enough data for such a fine-grained analysis. For the slider rating task, fixing a subject and an item, observations and replicates are probability vectors of length 15. In a first analysis, we used the mean of these probability vectors as a test statistic. The minimum posterior predictive p -value over all 20 (subjects) times 8 (items) cases was 0.13, suggesting that the means of observed s_{ij} are non-surprising to



(a) slider rating



(b) number estimation



(c) bin comparison

Figure 4: Posterior predictive checks for aggregate data. Red lines give empirical observations. Black lines are means of posterior predictive samples, gray areas are 95% HDIs.

the trained model. In a second analysis, we used entropy as a test statistic. Two cases gave posterior predictive p -values lower than 0.05. These were from the two participants who gave a very extreme slider rating for the “marbles” item, basically assigning all “mass” to the last bin. What this suggests is that the model can cope reasonably well also with individual-level data, but has problems accounting for “extreme” choices, given that the population-level hyper-prior on P_{ij} will lead to shrinkage. [ndg: can say that one possibility (future work) would be to switch to a fatter-tailed distribution relating Q to P , in order to deal with outliers.]

Conclusions

Acknowledgments

MF’s work was supported by the Institutional Strategy of the University of Tübingen (Deutsche Forschungsgemeinschaft, ZUK 63) and the Priority Program XPrag.de (DFG Schwerpunktprogramm 1727).

References

- Andersen, S., Fountain, J., Harrison, G. W., & Rutström, E. E. (2014). Estimating subjective probabilities. *Journal of Risk and Uncertainty*, 48, 207–229. doi: 10.1007/s11166-014-9194-z
- Degen, J., Tessler, M. H., & Goodman, N. D. (2015). Wonky worlds: Listeners revise world knowledge when utterances are odd. In *Proceedings of cogsci*.
- Gelman, A., & Rubin, D. B. (1992). Inference from iterative simulation using multiple sequences (with discussion). *Statistical Science*, 7, 457–472.
- Griffiths, T. L., & Tenenbaum, J. B. (2006). Optimal predictions in everyday cognition. *Psychological Science*, 17(9), 767–773.
- Kao, J. T., Wu, J. Y., Bergen, L., & Goodman, N. D. (2014). Nonliteral understanding of number words. *PNAS*, 111(33), 12002–12007. doi: 10.1073/pnas.1407479111
- Lee, M. D., & Wagenmakers, E.-J. (2015). *Bayesian cognitive modeling: A practical course*. Cambridge, MA: Cambridge University Press.
- Lewandowsky, S., Griffiths, T. L., & Kalish, M. L. (2009). The wisdom of individuals: Exploring people’s knowledge about everyday events using iterated learning. *Cognitive Science*, 33, 969–998. doi: 10.1111/j.1551-6709.2009.01045.x
- Manski, C. F. (2004). Measuring expectations. *Econometrica*, 72(5), 1329–1376. doi: 10.1111/j.1468-0262.2004.00537.x
- Morgan, M. G., & Henrion, M. (1990). *Uncertainty: A guide to dealing with uncertainty in quantitative risk and policy analysis*. Cambridge, UK: Cambridge University Press.
- Plummer, M. (2003). JAGS: A program for analysis of Bayesian graphical models using Gibbs sampling. In K. Hornik, F. Leisch, & A. Zeileis (Eds.), *Proceedings of the 3rd international workshop on distributed statistical computing*.

- Savage, L. J. (1971). Elicitation of personal probabilities and expectations. *Journal of the American Statistical Association*, 66(336), 783–801. Retrieved from <http://www.jstor.org/stable/2284229>
- Schlag, K. H., Tremewan, J., & van der Weele, J. J. (online first). A penny for your thoughts: a survey of methods for eliciting beliefs. *Experimental Economics*. doi: 10.1007/s10683-014-9416-x
- Schöller, A., & Franke, M. (2015). Semantic values as latent parameters: Surprising *few* & *many*. In S. D'Antonio, M. Moroney, & C. R. Little (Eds.), *Proceedings of SALT 45* (pp. 143–162). doi: 10.3765/salt.v25i0.3058